

NATS clustering: how far do three nodes actually **get you?**

Three demos in this lab were built to answer one growth question, and every claim on the slides that follow is a measured number, not a design opinion.

THE QUESTION

We start on one three-instance cluster. When must we add a second cluster, a gateway, an arbiter site or a leaf node — and what does JetStream cost us at each step?

THE EVIDENCE BEHIND THESE SLIDES

■ Demo 02 — multi-region · 6 servers, gateway, accounts

■ Demo 03 — 9 topologies · 178 checks, 0 failed

■ Demo 04 — JetStream + CQRS · rehydration cost

01

WHAT WAS ACTUALLY RUN

Three demos, one growth question

Each demo holds everything still except one variable. Demo 02 varies the account model on a fixed topology. Demo 03 varies the topology itself. Demo 04 drops to a single server and measures what JetStream costs when you rebuild state from the log.

DEMO 02 — MULTI-REGION

- > **6 servers**, two clusters of three, joined by one **gateway**.
- > Variable: the **account model**. One shared account, or one per region.
- > Answers: does a tenant wall hold across a gateway? What does **Replicas: 1** cost? Where does the data physically live?

DEMO 03 — TOPOLOGY AS THE VARIABLE

- > **9 runnable topologies** — five shapes and four variants — needing 6, 7 or 9 servers.
- > Every shape built from nothing by a script, measured, then torn down.
- > **178 checks passed, 0 failed**, plus 52 recorded observations. A check has an expected answer and passes only on an exact match.

DEMO 04 — JETSTREAM + CQRS

- > **1 server**. No cluster, no gateway, no domain.
- > Variable: rebuild an aggregate **with** a snapshot, or **without** one.
- > Answers: what does replay cost, and how fast does that cost grow?

THE FIVE WORDS THAT KEEP GETTING CONFUSED

WORD	QUESTION IT ANSWERS	PLAIN MEANING
Region	Where does it run?	A place. Cape Town. Sydney.
Tenant	Whose is it?	A paying customer. A commercial fact, not a running thing.
Account	What may it touch?	A NATS wall. It decides what you may see and send.
Cell	What ships together?	One stamp of the stack, deployed as a unit.
JetStream domain	Which storage pool holds the log?	Which NATS storage the messages actually live in.

WHY THIS SLIDE EXISTS

These five are **independent axes**. No two may be defined as permanently equal. One V3 drawing used the word **“domain” 26 times** on one page, carrying three unrelated meanings.

Only one is enforced by a machine: **NATS enforces the account wall**. The rest are conventions we hold ourselves. Region and cell are places; tenant and account are not.

02

FIVE WORDS USED ON EVERY SLIDE THAT FOLLOWS

Instance, cluster, gateway, leaf — and JetStream namespace

The shape below is the one we would actually deploy: two AWS regions, three availability zones each, one NATS server per zone. Three zones is why a cluster is three servers — lose a zone and two of three still holds a majority.



1 · NATS INSTANCE

One server process. It holds data and answers clients. Here, one per availability zone.

2 · CLUSTER

Three or more instances that **vote together**, inside one region. This is the unit of availability, and the only thing three servers buys you.

3 · GATEWAY

Joins two clusters and routes messages between them. It does **not** copy data — and it fuses both sides into one **meta group**: the set of servers that vote on every stream change. See slide 06.

4 · LEAF NODE

A server that borrows a hub over one outbound link. It may run its own JetStream. A leaf link is the only place `jetstream.domain` may change. See slide 10.

5 · JETSTREAM NAMESPACE

The stream and KV names **one meta group** manages. **Not the account**, which owns the subjects. A name is unique per `meta group × account` — change either one and the collision goes away.


03

THE THREE-INSTANCE CLUSTER

Three servers is not three copies

A cluster is three instances because RAFT needs a majority to elect a leader, and the majority of three is two — so one instance can be lost. That protects the *cluster*. It does not protect your *stream*. Replication is a separate, per-stream choice, and its default is **one**.

REPLICAS: 1 — DEFAULT



stop the one server holding it

stream is offline (10118)

DATA STUCK UNTIL THAT SERVER RETURNS

REPLICAS: 3 — ASKED FOR, PER STREAM



stop the stream-leader

new stream-leader in about one second

NOTHING LOST · EVERY WRITE NOW COSTS THREE SERVERS

THE TRAP

- › The default is **1**, and nothing warns you. A three-server cluster with an R1 stream survives nothing.
- › **Replicas never cross a region.** `--replicas 3` means three copies in **za** — not one in each region.
- › R3 is slower and uses three times the disk. Choose it **on purpose, per stream**.

WHAT `--REPLICAS 3` ACTUALLY COPIESCOPIED 3 TIMES, INSIDE **ZA**

- › **Stream messages** — the log itself.
- › **Stream config** — subjects, limits, retention.
- › **Consumer position and acks** — its *own* vote group, which inherits 3 from the stream.
- › **KV bucket values** — a KV bucket **is** a stream.
- › **Object store blobs** — also a stream.

HANDLED SOMEWHERE ELSE

- › The **fact** that the stream exists. That lives in the meta group, on every server in the meta group — not in the replica set.

NOT COPIED AT ALL

- › **Anything in **au**** . A gateway routes; it does not copy. That needs a mirror.
- › **Ephemeral consumers** — R1, in memory, and they die with their server.

RECOMMENDATION


Start at three instances, one cluster. Set **Replicas: 3** explicitly on every stream and KV bucket that matters. Make it a reviewed default, not a habit.

04

FIVE WAYS TO ADD A SECOND REGION

Every shape is majority arithmetic

Demo 03 built nine topologies from nothing and measured each one. The single question that separates them is: **how many servers are in the JetStream meta group, and who still holds a majority when a region goes dark?** This slide only introduces them. The findings follow, one shape at a time.



Red ring = one shared fate. Green ring = one shared fate that still holds a majority when a region is lost. Amber rings = separate meta groups, each with its own vote.

THE FOUR VARIANTS

VARIANT	WHAT IT ASKS
B · domain per cluster	T2, plus a different <code>jetstream.domain</code> on each side. Can a domain name split a supercluster?
E · export / import	T2, plus two accounts instead of one. Can a region share data deliberately, across the account wall?
G · T6	A gateway and a hub leaf link at once. Does the leaf link change what the gateway does?
H · T5 + accounts	T5 with an account per region. Does that remove double capture?

WHAT EACH SHAPE IS

SHAPE	SERVICES	META GROUPS	WHAT IT IS
T1 islands	6	2	Two clusters, no link at all. Each region is its own NATS and its own JetStream.
T2 gateway	6	1	The same two clusters, joined by one gateway.
T3 + 1 arbiter	7	1	T2 plus a single-node third site that also votes.
T4 + 3 arbiters	9	1	T3 with a full three-node third site instead of one node.
T5 hub + leaf	9	3	A hub cluster, with a leaf cluster in each region. Each runs its own JetStream.

WHAT FOLLOWS


One finding set per shape: **05** T1 · **06–07** T2 · **08** B · **09** T3 and T4 · **10** T5 · **11** H · **12** E · **13** G. Every row on them is a check with an expected answer that matched.

05

T1 · NO LINK AT ALL · 6 SERVERS · 12 CHECKS, 12 PASSED

Two islands never freeze, and never agree

The baseline. Two clusters of three, with nothing joining them. Each region runs its own JetStream and its own vote. This is the only shape where a region going dark changes nothing at all for the other one — and the only shape that shares nothing whatsoever.



WHAT IT GIVES YOU

- Two separate meta groups, of three servers each. A majority of three is two, so each region tolerates losing one server. T1a , T1b , T1c
- ZA goes dark and AU does not notice. AU keeps its meta leader, creates a brand-new 3-replica stream, and its meta group size is unchanged at 3. T1g , T1h , T1i
- No WAN on any read or write. Nothing to fail between the regions, because there is nothing between them.

WHAT IT COSTS YOU

- Two streams can carry the same name and be unrelated. AU asks for ODOMETER , gets it, and it is a second real stream. No error, because there is no shared JetStream namespace to collide in. T1d , T1e
- Nothing crosses. One publish in ZA: ZA holds 1 message, AU holds 0. T1f
- Every cross-region copy is an application you have to write, run and monitor yourself.

MEASURED — DEMO 03, T1			
CHECK	QUESTION	EXPECTED	MEASURED
T1a	Distinct JetStream meta groups across both regions	2	2
T1f	One publish in ZA — messages in ZA / AU	1 / 0	1 / 0
T1g	ZA dark — AU still has a meta leader	yes	yes

CHOOSE T1 WHEN

The two regions genuinely share no data, and you would rather write one explicit copier than accept a shared vote. It is the safest shape for availability and the worst for consistency — and those are the same fact.


06

THE MOST EXPENSIVE THING ON THESE SLIDES

A gateway makes one JetStream meta group, not two

A gateway does not link two meta groups — it fuses them into one supercluster with **one meta group** and **one leader election**. The account wall still holds for data. It does not hold for availability.

WAN UP · ONE META GROUP · SIZE 6 · MAJORITY 4



meta leader: za-2 · 6 OF 6 · ONE VOTE, WRITES EVERYWHERE

ONE META GROUP · SIZE 6 · MAJORITY 4 · INFERRED — A LIVE CUT WAS NOT RUN



WAN cut (inferred)

meta leader: NONE

3 OF 6 — NOT A MAJORITY

meta leader: NONE

3 OF 6 — NOT A MAJORITY

READ THE TWO DRAWINGS TOGETHER

Only the **link** changed — and with it, the arithmetic. Six votes become three, three is not four, and every JetStream **change** stops in both regions at once. Measured: a whole region **stopped**. A live two-sided cut is inferred.

07

DEMO 03 · T2 / A · 6 SERVERS · 26 CHECKS — ALL PASSED

One JetStream namespace, one copy, and one shared fate

The previous slide showed the arithmetic. These are the consequences you will actually meet: losing a whole region is the same sum as losing the link, names collide, data lives in exactly one place, an uneven split behaves differently — and some things carry on regardless.

IF A WHOLE REGION DIES, NOT JUST THE LINK

- › The meta leader's **location** does not matter. Lose any 3 of 6 and you are under 4. [A10](#)
- › AU still publishes into the streams **it** holds. It cannot create anything new. [A11](#) , [A12](#)
- › ZA's streams are not in AU. A gateway **routes**; it never copies. [A8](#)
- › [/jsz](#) went on naming the dead meta leader — **for up to about a minute**. [A10a](#)

ONE JETSTREAM NAMESPACE MEANS ONE NAME

- › AU asks for a stream ZA already owns → **10058 stream name already in use**
- › Overlapping subjects → **10065 subjects overlap**
- › So inside one account, two streams cannot **capture** the same subject. A mirror, a source or a repeated publish can still make a second copy.
- › **Different accounts, no collision.** [LB_ZA](#) and [LB_AU](#) may each own an [ODOMETER](#) — the wall is the **account**, not the gateway. Slide 11, [E7](#) , [E8](#)

...AND ONE COPY, IN ONE PLACE

- › AU's client keeps working, because a stream is reachable from either side of a gateway.
- › But [stream info](#) from AU reports cluster **za**. AU holds nothing.
- › AU pays the WAN on every read and goes dark for that stream when ZA is lost.

UNEVEN REGIONS BEHAVE DIFFERENTLY

- › A **3 / 3** split freezes both sides.
- › A **3 / 2** split does not — the side holding 3 of 5 keeps its leader.
- › Which side survives is arithmetic, so design the count deliberately.

WHAT A 3/3 SPLIT FREEZES

- › [stream add / rm / edit](#) → **10008 JetStream system temporarily unavailable**
- › Every account at the same moment. Measured with 3 JetStream accounts live: **1 meta group**, not three.

WHAT KEEPS WORKING

- › Core NATS publish and subscribe.
- › Reads from a stream **your own region holds**. [A20](#)
- › **Writes into a stream that already exists** — measured 10 → 15 messages, all acked.
- › **Not** a read of a stream the dark region held — that fails **10008** . [A21](#) , [A21a](#)

THE DECISION

If both regions must accept a JetStream **change** while the other is dark, a plain two-region gateway cannot do it. Add a vote outside both regions, or stop sharing a meta group.

08

THE SETTING EVERYBODY REACHES FOR FIRST

A domain may only change across a leaf link


`jetstream.domain` looks like the fix for the previous slide. It is not. We set a **different** domain on each side of a gateway, and it made things quietly worse. One domain across a whole supercluster is supported — it just splits nothing. Over a **leaf** link a domain per side is real.

A · DOMAIN OVER A GATEWAY — DOES NOT WORK



meta group size still 6 — domains did NOT split it
 one side never elects · placement → 10005 no suitable peers
 AND WHICH SIDE GOES BLIND IS A COIN TOSS — MEASURED BOTH WAYS

B · DOMAIN OVER A LEAF LINK — REAL



OVER A GATEWAY — MEASURED

- Meta group size **still 6**. A domain name cannot split a supercluster.
- One cluster elects, the other reports `leader: null` forever.
- Placing a stream on the blind cluster → `10005 no suitable peers for placement`
- The verdict in the report is blunt: **does not work**.

OVER A LEAF LINK — MEASURED

- 3 distinct meta groups**, sizes 3 / 3 / 3. Hub, ZA and AU each own their JetStream.
- The whole hub can be stopped and both regions still elect, still create R3 streams, still store writes.
- What sets T5 apart is an **independent quorum per region**. T6 also survives hub loss (G16 – G18), but its regions still share one vote.

THE PRICE NOBODY MENTIONS

- Double capture is real here** on a shared account. One publish in ZA was stored **1 and 1** — twice, with no warning. An account per region removes it (slide 11).
- Nothing replicates by itself**. Every cross-region copy is hand-written.
- A mirror over a leaf needs `mirror.external.api` → `$JS.au.API`. Without it: **0 messages copied**, and the stream reports healthy.

THE RULE



Use one domain across a supercluster, or none. A different domain per gateway-joined cluster does not split the meta group — it leaves one cluster blind while the config reports success. Change the domain only across a leaf link.

09

T3 · 7 SERVERS · 18 CHECKS — T4 · 9 SERVERS · 26 CHECKS — ALL PASSED

Buy a vote that sits outside both regions

T2 freezes because 3 of 6 is not a majority. Both fixes put more voters in a third site, so a surviving region plus that site still adds up. T3 adds one node, T4 adds three, and the difference is **slack**.



T3 — ONE ARBITER NODE · 7 SERVERS, MAJORITY 4

- One meta group across all three sites. [C1](#) , [C2](#)
- Either region can go dark and the survivors elect a leader and still create a new 3-replica stream. [C3](#) — [C6](#)
- Zero slack. 4 of 7 is exactly the majority. One more node after the region loss and it freezes. [C12](#)

T4 — A THREE-NODE ARBITER SITE · 9 SERVERS, MAJORITY 5

- Still one meta group, size 9. [F1](#) , [F2](#)
- A region dark leaves 6 of 9 — a leader is elected and writes continue. [F3](#) , [F4](#)
- One node of slack. A region *plus* one arbiter node leaves 5 of 9, exactly the majority, and it still works. At 4 of 9 the leader is NONE. [F5](#) , [F6](#) , [F7](#) , [F11](#)
- Even with no meta majority, publishing into an **existing** stream still works — provided that stream keeps its own quorum. Only changes stop. [F8](#) , [F9](#) , [F21](#)

THE ARBITER TRAP — IT WILL QUIETLY HOST YOUR DATA

- An explicit `--cluster arb` R1 stream **lands on the arbiter**, and a client dialling the arbiter's own port creates streams there too. [C8](#) , [C10](#)
- Asking for R3 on a one-node site fails with [10005](#) . [C9](#)
- So: keep clients off the arbiter's client port, or fence it with placement. [C11](#)

ONE NUMBER WORTH REMEMBERING

After the majority was gone, `/jsz` went on naming a meta leader for **tens of seconds**. [F7a](#)

And a new leader is **named before it will take a change** — ask at once and a healthy supercluster refuses. [F6a](#)

Late in one direction, early in the other. Neither edge survives a single reading.

A WHOLE REGION DIES — AND THE ANSWER FLIPS


- ZA dies holding the meta leader. **T3 keeps 4 of 7, T4 keeps 6 of 9** — still a majority, so the survivors elect one and AU creates **new** streams. [C5a](#) , [C6](#) , [F4](#)
- The leader's **address** never mattered. Only the count does.
- ZA's data is still gone**, even with the vote alive. [F28](#)
- On slide 06 the same loss leaves 3 of 6 and freezes both. **That is what the third site buys.** [A10](#)

10

T5 · 9 SERVERS · 26 CHECKS — ALL PASSED

Independence you can have, at the price of copying by hand

With this lab's leaf link (`$SYS` not carried, a domain per side), each region keeps its own JetStream and vote. The hub can vanish and both regions carry on. Nothing replicates by itself.



T5 — HUB AND LEAF · THREE SEPARATE META GROUPS

- › **Three meta groups**, sizes 3 / 3 / 3. Each has its own quorum of two. [D1](#) , [D2](#)
- › **The whole hub can vanish.** Both regions keep a leader, both create new 3-replica streams, and AU keeps storing its own publishes. [D10](#) – [D14](#)
- › Both regions may own a stream called `ODOMETER` with no `10058`, because there is no shared JetStream namespace. [D3](#) , [D4](#)

T5 — WHAT IT COSTS

- › **Double capture is real here.** ONE publish in ZA was stored in ZA *and* in AU — 1 / 1 — and two same-named shadow streams lived at once. [D5](#) , [D6](#)
- › **That 1 / 1 is not the gateway's 1 / 1** — there it is one stored copy read twice. An account per region makes it 1 / 0; next slide. [D5a](#) , [D18](#) , [H6](#)
- › **Nothing replicates by itself.** Every cross-region copy is hand-written, and the panel below is why that is harder than it sounds. [D15](#)

SO HOW DOES DATA GET FROM ONE LEAF TO THE OTHER? YOU WRITE IT BY HAND.

- › A leaf link carries **messages**, not the catalogue. A name-only lookup stays local. [D1](#)
- › What works is a **mirror naming the remote API**: `external.api $JS.au.API` → 2 copied. [D8](#)
- › **Leave that line out and you get silence.** 0 messages, no error, no warning, stream healthy. [D7](#) , [D7a](#)
- › **Worse:** a name-only mirror silently copies the *local* stream of that name. [D9](#) , [D9a](#)

Rule: always write `external.api`. A green mirror is not evidence that anything was copied.

11

T5 / H · 9 SERVERS · 15 CHECKS — ALL PASSED · MEASURED 2026-09-23

The fix for double capture, and the bill it hands you

The previous slide named double capture as T5's price, and per-region accounts as the way out. That is now measured, on the same nine servers. It works — and it takes the hub down with it.

WHAT THE ACCOUNT WALL BUYS

- › **Double capture is gone.** One publish in `LB_ZA` → 1 in ZA, 0 in AU. On the shared account the same send was 1 / 1. `H6` , `D18` , `D18a`
- › **Still three meta groups**, 3 / 3 / 3, quorum two each. The account changes nothing about the vote. `H1` , `H2`
- › **The hub is a reader, not a dependency.** Kill all three hub servers and a region keeps its leader, creates new 3-replica streams, and stores what it accepts. `H12` – `H16` , `H14a`
- › **A whole region can go dark** and the other carries on, then both come back to 3 / 3 / 3. Measured on T5's shared account (`D24` – `D27`); inferred for H.

WHAT IT COSTS

- › **The leaf link binds ONE account.** An account with no remote of its own is **six separate islands**, one per site (T5 / D). `D22a`
- › **H gives each regional account its own remote.** The hub then captures and mirrors each region — it removes ZA / AU cross-capture, not capture at the hub. `H8` , `H11`
- › **A shared hub is not a shared account.** Both regions reach the same three hub servers and still cannot see each other. `H6a` , `H9` , `H9a`
- › **The mirror trap is unchanged.** Name-only → 0 copied; with `external.api` → 2. The account was never the problem — `external.api` names a **domain**. `H10` , `H11` , `H11a`

THE SAME TWO NUMBERS, THREE DIFFERENT MEANINGS — READ THEM WITH THE SHAPE IN HAND

- › **Gateway, shared account — 1 / 1.** ONE stored copy, read from two sides. `A26` , `C13` , `F13`
- › **Leaf link, shared account — 1 / 1.** TWO stored copies of one send, with their own sequence numbers. `D5` , `D5a`
- › **Leaf link, account per region — 1 / 0.** One copy, and the other side never saw the subject. `H6`
- › **So a passing check is not a finding.** The same measurement means the opposite thing on the shape above it.

THE DECISION


If double capture is what frightens you, an account per region on a leaf link removes it outright. Accept in exchange that a name-only request from the hub resolves in the hub's own domain. To read a region's stream it must name that domain's API, as the `H11` mirror does — and every copy is still hand-written. `H7a` , `H8a` , `H15`

12

E · 6 SERVERS · 17 CHECKS — ALL PASSED · THE ONLY SHARING IN THIS DEMO THAT IS EXPLICIT

The one safe way to copy across a wall

E is **not** T5. It has two clusters and a gateway, and no leaf nodes at all. It answers a different question: if two **accounts** must share one subject, how do you do it so nobody can ever confuse the copy with the original?



E · export / import

TWO CLUSTERS AND A GATEWAY · NO LEAF NODES

6 SRV · BOTH ACCOUNTS ON BOTH CLUSTERS · ONE WAY

WHAT IT DOES

- › ZA publishes once. The owning account stored it, and the importing account **LB_AU** stored it too, across the gateway. **E1** , **E2**
- › **It arrives renamed.** **LB_AU** 's raw-subject stream took **none** of it — the importing side prefixes the subject. **E3**
- › So imported traffic sits on subjects of its own, apart from local traffic. **E3a** That is **not proof of origin**: a local user can publish onto the prefix too, unless permissions reserve it. The rig set none.
- › When **LB_AU** publishes the same subject itself, ZA's stream is **unchanged**. Traffic only moves the way you declared. **E4** , **E5**

THE ACCOUNT WALL STILL HOLDS

- › **LB_AU** asking for **stream info** on ZA's **ODOMETER** **fails**. **E6**
- › And **LB_AU** may create its **own** stream called **ODOMETER** with **no 10058**. **E7**
- › They are two real, separate streams, in **za** and **au**. **E8**
- › Compare slide 06: one account over a gateway is **one** JetStream namespace, and that name collides.

SHARE THIS WAY

Explicit, one direction, and **renamed on arrival** — the only sharing measured in this demo that a reader cannot misread. **E9** Wherever two tenants or two regions must see the same event, put an account wall between them and export across it. Do not share by giving two streams the same name and hoping.


demo 03 · E1–E18 · ADR-054, ADR-056

13

ONE SERVER → ONE CLUSTER → TWO REGIONS · AND THE ONE STEP THAT IS NOT A STEP

Growing is cheap until the last move

Every other slide shows a **shape**. This one shows the **order**. Three of the moves are small and ordinary. The last one is not a small move. The lab measured two in-place shortcuts failing ([G24](#) , [G19](#) – [G22](#)). A conversion that keeps existing stores was never run ([D03-R11](#) , [D03-R12](#) , unmeasured). The cutover itself is a proposal.



The blue arrows are **design, not measurement** — demo 03 built every shape from nothing and never ran a migration. What is measured is [G24](#) : a leaf link added beside a gateway changes nothing the gateway does.

STEPS 1 TO 3 — ORDINARY

- One server is a starting point, not a design. It has no quorum and no HA.
- Three servers in one cluster is the first real shape: majority 2, one node may die.
- Going from one server to a cluster, **every** member needs at least one route — the seed included. A routeless seed with JetStream on refuses to start.
- Step 3 is where the whole decision lives — and it is a fork, not a line.

WHY THE SHORTCUT FAILS — MEASURED

- T6 put a gateway **and** a hub leaf link on the same servers. Result: **two** meta groups, not three. [G1](#) , [G2](#)
- The gateway still fuses both regions into one group of 6, and both name the same meta leader. [G3](#) , [G15](#)
- [10058](#) comes back, and T5's double capture collapses to **one** stored copy read through two handles. [G4](#) , [G7](#) , [G7a](#)
- A dark region still freezes every create with [10008](#) , and a healthy hub cannot lend a vote. [G11a](#) , [G14a](#)

SO DECIDE ON DAY ONE

- If a region must survive alone**, go 1 → 3 → hub and leaf. Do not stop at a gateway on the way.
- If one JetStream namespace is the point**, the gateway is the simpler shape — then buy the vote back with a **three**-node arbiter site. One node of slack, not zero.
- Changing your mind later is a **proposed** parallel build and cutover: stand up T5 beside T2, copy, switch clients, remove the gateway. It is not yet validated.

PICK THE END SHAPE FIRST

Adding the hub link early buys nothing — the gateway decides for as long as it exists. [G24](#) Whether any gradual path exists is **not measured**. The cheap moves are 1 → 3 → a second region; the expensive move is changing which kind of second region you chose.

demo 03 · T6 / G · G1–G24 · the migrations themselves are unmeasured

14 WHAT THIS LETS SOMEBODY DECIDE

Pick the shape from the requirement, not the diagram

Every row below is a requirement somebody will actually state, and the shape the measurements say it selects. The mistakes happen when a shape is chosen first and a requirement is fitted to it afterwards.

IF THE REQUIREMENT IS...	...THE SHAPE IS	...AND THIS IS WHAT IT COSTS
One tenant must not see another tenant's data	An account. Not a cluster, not a domain.	Nothing. It is the one wall NATS enforces itself, and it holds across a gateway.
One name, one copy, no overlapping capture	Gateway (T2)	Shared availability. A 3/3 split freezes every account at once.
Both regions must accept a JetStream change while the other is dark, in one shared JetStream namespace	A vote outside both regions — T3 (7) or T4 (9). Without the shared namespace, T5 does this too.	A third site. T3 has zero slack afterwards; T4 has exactly one node .
A region must keep running when the centre is gone	Leaf links + a domain per region (T5)	Double capture on a shared account (an account per region removes it), and every cross-region copy hand-written as a mirror.
Two accounts must share one subject on purpose	Export / import	One-directional, read-only, renamed on arrival. The only safe sharing measured.
Data must stay in-country (residency)	A regional cell , gateway-joined, under the one global operator	Placement discipline: the tenant is pinned, no stream mirrors or sources leave the cell.
Access and control must stay in-country (sovereignty)	A sovereign cell — own operator, own resolver, no gateway	A second provisioning act, a second key chain to rotate, a second resolver to run.
A depot or yard must survive losing its link	A leaf node	Not a cross-border answer by our policy. NATS lets a leaf bridge operators and authenticate locally; the rule is ours, not the server's.

WHERE TO START — AND STAY

- > **One cluster of three.** It survives one loss, and it is the whole of what most of the platform needs.
- > **Replicas: 3** set explicitly on every stream and KV bucket that matters.
- > Accounts as the tenant wall from day one. Retrofitting that boundary later is the expensive move.

WHAT WE HAVE NOT MEASURED YET

- > The **latency** cost of a cross-region read. Placement is measured; latency is not.
- > **Mirror lag and bandwidth** in either shape — so the delay a mirror adds to a CQRS read model is unpriced. Slide 16.
- > The hub as the **primary store** for both regions under load, and what losing it costs then. It does hold data of its own (**H8** , **H11**); load and loss are unmeasured.
- > Leaf reconnect with more than two regions.

THE NEXT QUESTION

Every shape here is available today and we can pick one on the evidence. What we cannot yet price is **the wire**: what a cross-border byte costs in latency and in bandwidth.

That is the next demo, and it is the one that decides whether a gateway between regions is affordable at all.

15 WHERE EVERY NUMBER CAME FROM

A figure with no provenance becomes a stale constant

Each number on these slides, the day it was measured, and what it ran on. Nothing here was taken from documentation or from memory.

THE RUNS

SOURCE	MEASURED	ON	WHAT IT PRODUCED
Demo 03 — <code>lab/run-all.sh</code>	2026-09-28 13:40 SAST	nats-server 2.14.6 · Darwin 25.4.0 arm64	178 checks passed, 0 failed, 52 notes, across 9 topologies. Meta group sizes, 10008 / 10058 / 10005 refusals, mirror behaviour, all six findings.
Demo 02 — multi-region lab scripts	2026-09-09 & 2026-09-11	nats-server 2.14 in Docker, one host	Replicas 1 vs 3, the account wall across a gateway, WAN cut = change freeze only (10 → 15 messages acked), subject overlap 10065.
Demo 04 — <code>cqrs rehydrate</code>	2026-09-16	nats-server 2.14.3 in Docker, on a laptop	10 001 events on one vehicle, best of three runs each way: 25 ms without a snapshot, 1.1 ms with one.
ADR-054 / ADR-056	Accepted 2026-09-07 / 2026-09-18	V3 platform architecture record	The five axes, the tenant-wall rule, and the three cross-border shapes (leaf node, regional cell, sovereign cell).

RETRACTIONS — BOTH CORRECTIONS ARE ON THE RECORD

- › “25 / 25 / 50 km cross-region double capture” is withdrawn. The 25 came from a durable pull consumer being deleted on every projector run, so each run replayed the whole stream. Four runs over a stream holding one message read 12.5, 25, 37.5, 50. The 50 came from adding one region’s number to the other while both read the same bucket. Behind a gateway there is nothing to capture twice — double capture is real in the hub-and-leaf shape instead.
- › “8.2 s, about 600x” is withdrawn. The replay used an ordered consumer’s `Next()`, which resets the consumer on every call — the demo was creating one consumer per event. About 99% of that time was consumer bookkeeping. The corrected figure is 25 ms vs 1.1 ms, about 23x. The lesson survived the correction; the number now measures the lesson instead of a bug.

HOW TO REPRODUCE ANY OF IT


- › Demo 03, all nine topologies: `demos/03-multi-cluster-and-accounts/lab/run-all.sh` — builds every shape from nothing and regenerates `REPORT.md` and `REPORT.html`. Both are generated; neither is hand-edited.
- › Demo 02, replicas and the account wall: `demos/02-multi-region/lab/01-the-wall.sh`, `02-replicas.sh`, `03-odometer.sh`.
- › Demo 04, rehydration: `cqrs seed -vehicle V1 -n 10000`, then `cqrs rehydrate -snapshot=false`; wait about 20 seconds for the asynchronous snapshotter, then `cqrs rehydrate -snapshot=true`.
- › A caution that applies to all three: never fake a WAN cut with `docker network disconnect` — it drops published host ports, and an early run wrongly reported that both regions had died.

16 STREAMS, CONSUMERS, KV AND THE COST OF REPLAY

A KV bucket is a stream, and a consumer lives where its stream lives

Placement is not a separate subject from clustering — it is the same subject seen from inside one cluster. And once state is rebuilt from a log, the length of that log becomes a latency number you have to manage.

THE CQRS SHAPE MEASURED IN DEMO 04



The log is the only source of truth. Delete both KV buckets, reset the projector consumers to the start, and a replay rebuilds them. Delete the stream and nothing does.

REHYDRATION — 10 001 EVENTS ON ONE VEHICLE



Replay cost grows with the length of the log. The snapshot read is one KV get plus whatever tail is left, so its cost follows the **tail**, not the log. A tenfold longer log should cost roughly tenfold more to replay — an extrapolation, not a measurement.

PLACEMENT — THE RULES WE MEASURED

- › Placement is **per stream**. A new stream lands in the cluster of the client that creates it, unless you name one with `--cluster`. The account and the meta leader do not pin it.
- › A **KV bucket follows the same rule, because a bucket is a stream**.
- › A durable pull consumer created from AU for a ZA-held stream reported cluster **za**. AU held no part of it and paid the WAN on every pull.

Not measured. A local mirror takes the WAN off reads but adds its lag to every read model built on it. A cut-off mirror serves old data with no error. Check commands against the source, never a mirror.

TWO SETTINGS THAT FAIL SILENTLY

- › **LimitsPolicy, not InterestPolicy**, on any stream you intend to replay. InterestPolicy discards a message once every consumer has acked it, and the replay then returns nothing — with no error.
- › **The snapshot may be stale.** Rehydration must read the snapshot *and then replay the tail*. Measured mid-catch-up: 2 900 events of tail in 12 ms.

STORAGE — THE STANDING RULE

- › **A count is never shown without its bytes.** 100 000 000 events sounds reasonable right up to the moment you learn it is about **7.6 GB** — extrapolated from about 76 bytes per event.
- › R3 triples that figure, per stream, per region.
- › Seeding one million events into the bench fixture took **2.195 seconds**, measured — the log grows faster than anyone plans for.

RECOMMENDATION

Treat snapshotting as a **requirement**, not an optimisation, on any aggregate that will take more than a few thousand events. And state a retention policy on day one, in bytes.

17

FINDINGS THIS LAB ADDED · THE SIXTH ONE IS SLIDE 12

Five failures that report success

These are not in the documentation. Each one was found by a script that expected a specific answer and got a different one. Every one of them looks healthy from the outside — which is exactly why they are on a slide.

01 · THE MONITORING ENDPOINT LIES WHILE IT IS STALE

After the majority of a supercluster was frozen, `/jsz` kept naming a meta leader for **roughly a minute**. Read it too early and a dead supercluster looks perfectly healthy.

In that window the client does not get a clean refusal either — **it hangs and times out**. The clean `10008` only appears once the old leader has aged out.

02 · A NEW LEADER IS NAMED BEFORE IT CAN TAKE A CHANGE

After an election, `/jsz` names the new leader **seconds before** that leader will accept a create. Ask immediately and the call is refused — on a supercluster that is in fact healthy.

So **neither edge of a failover can be trusted from a single reading**. Health checks must retry and record how long they waited, not sleep a magic number.

03 · A NAME-ONLY MIRROR CAN COPY THE WRONG STREAM

A mirror with no `external.api` does not fail when the name it wants only exists in **another domain** (this lab's leaf link) — **it just sits at zero**. Across a gateway there is one domain, and a name alone works (`A15`, `A16`). Worse: if a stream of the same name also exists locally, the mirror copies *that* one, reports healthy, and looks like it is working.

Proved by giving the two streams different message counts. No error, no warning, wrong data.

04 · WHICH REGION GOES BLIND IS A COIN TOSS

Two runs of the same configuration, minutes apart, nothing changed — and the blind side swapped.

A team that tests this shape once will write down the wrong region's name. Our check deliberately refuses to name a side.

05 · A TEST THAT DOES NOT TEST READS EXACTLY LIKE A RESULT

Not a NATS fact — a rig fact, and the most important one. An early version of the harness recorded the wrong process id, so `kill -STOP` hit nothing and the “frozen” region stayed fully alive. **The lab reported healthy behaviour and it looked like a finding**. Every freeze in these scripts now proves itself first, by checking the monitor port has really gone dark.

TWO NUMBERS WE RETRACTED

- › A cross-region “double capture” of **50 km** was a replay bug — a durable consumer was being deleted on every run.
- › A rehydration speed-up of **600x** was **99% consumer bookkeeping**, not log reading. The real figure is 23x.
- › **A lab whose numbers only ever improve is not measuring**.